
presto

1. BIG DATA


HIVE
OOZIE
STORM
hadoop
FLUME
cloudera
APACHE
HBASE
Hortonworks
Spark
Cassandra

1.1. HADOOP

ARCHITECTURE

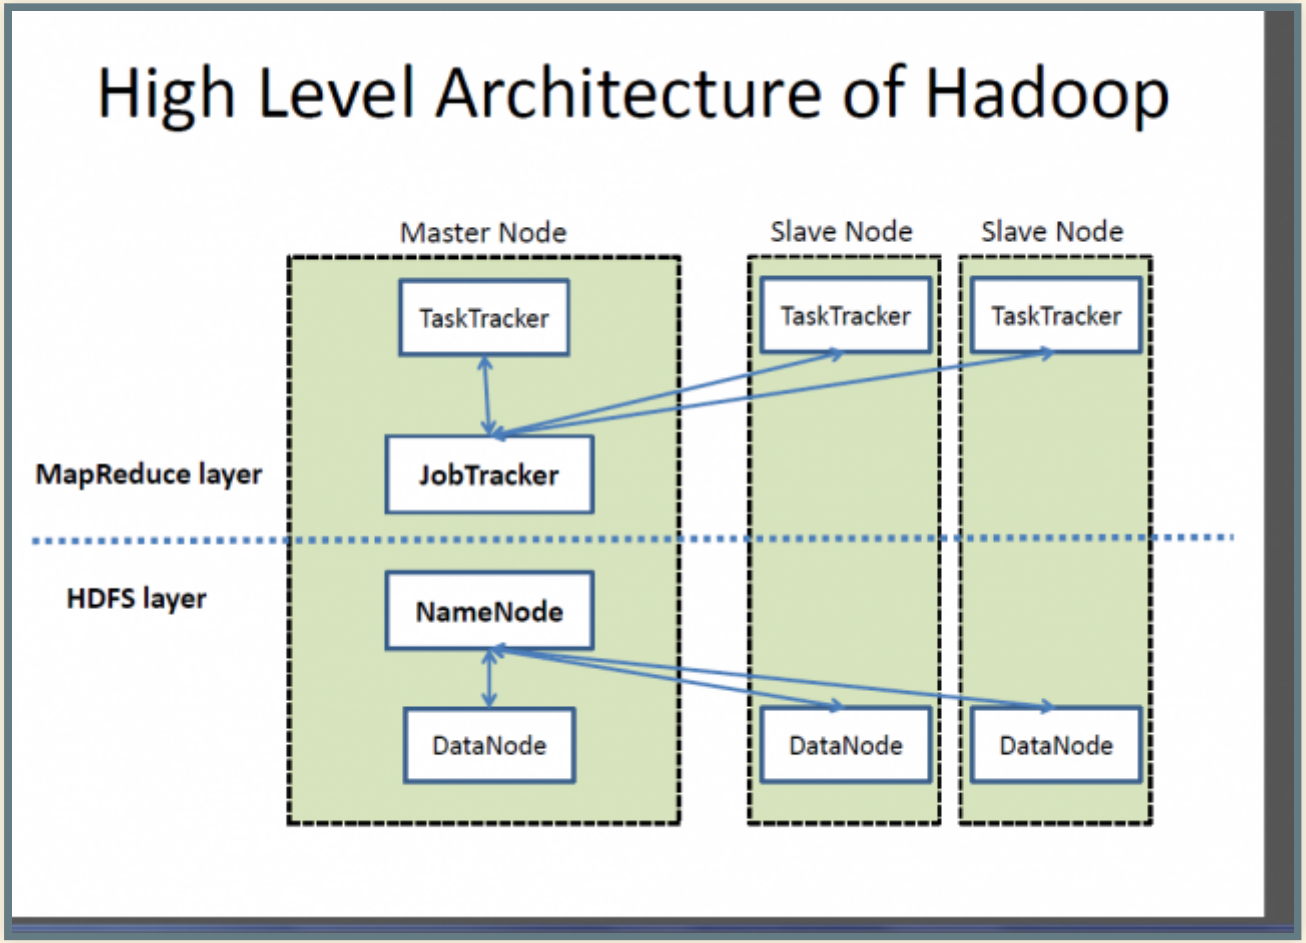

HDFS

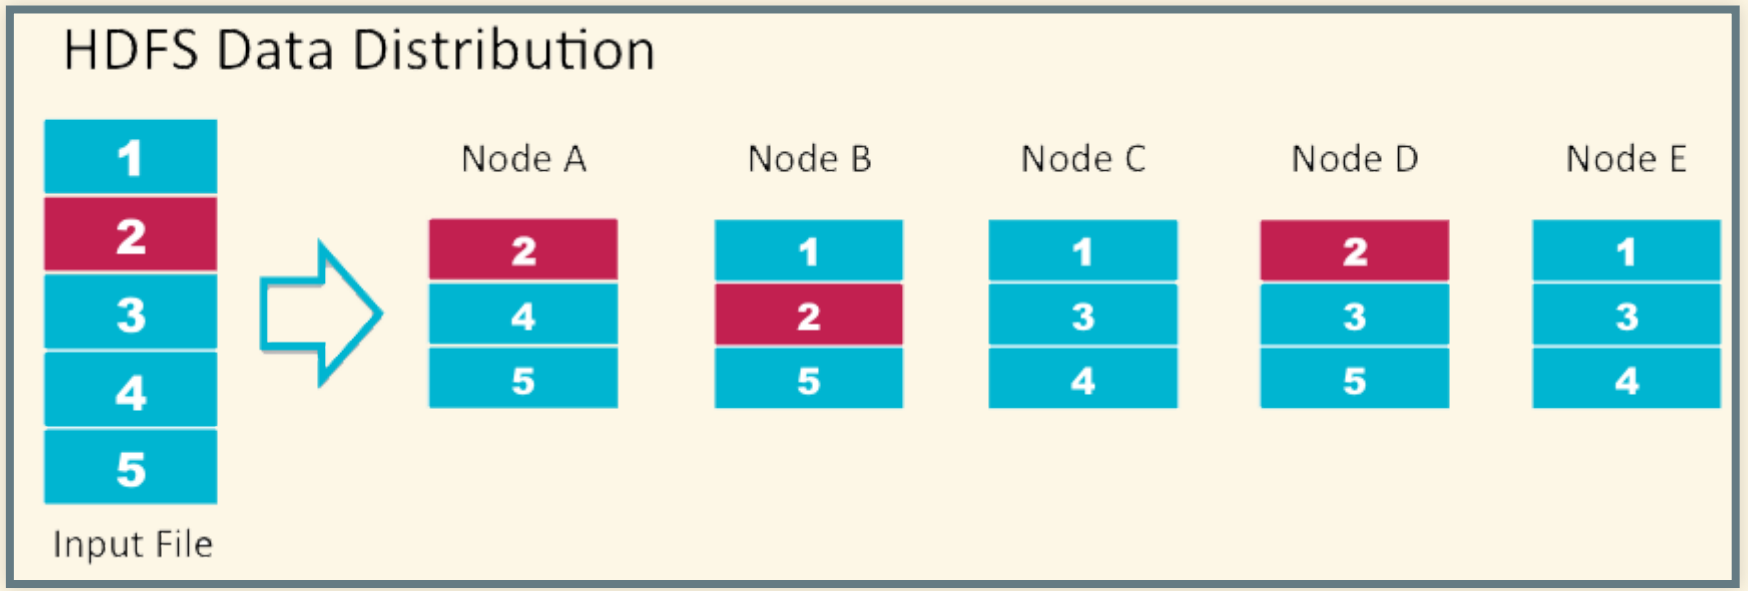
HDFS Data Distribution
1
2
3
4
5
Input File
Node A
2
4
5
Node B
1
2
5
Node C
1
3
4
Node D
2
3
5
Node E
1
3
4

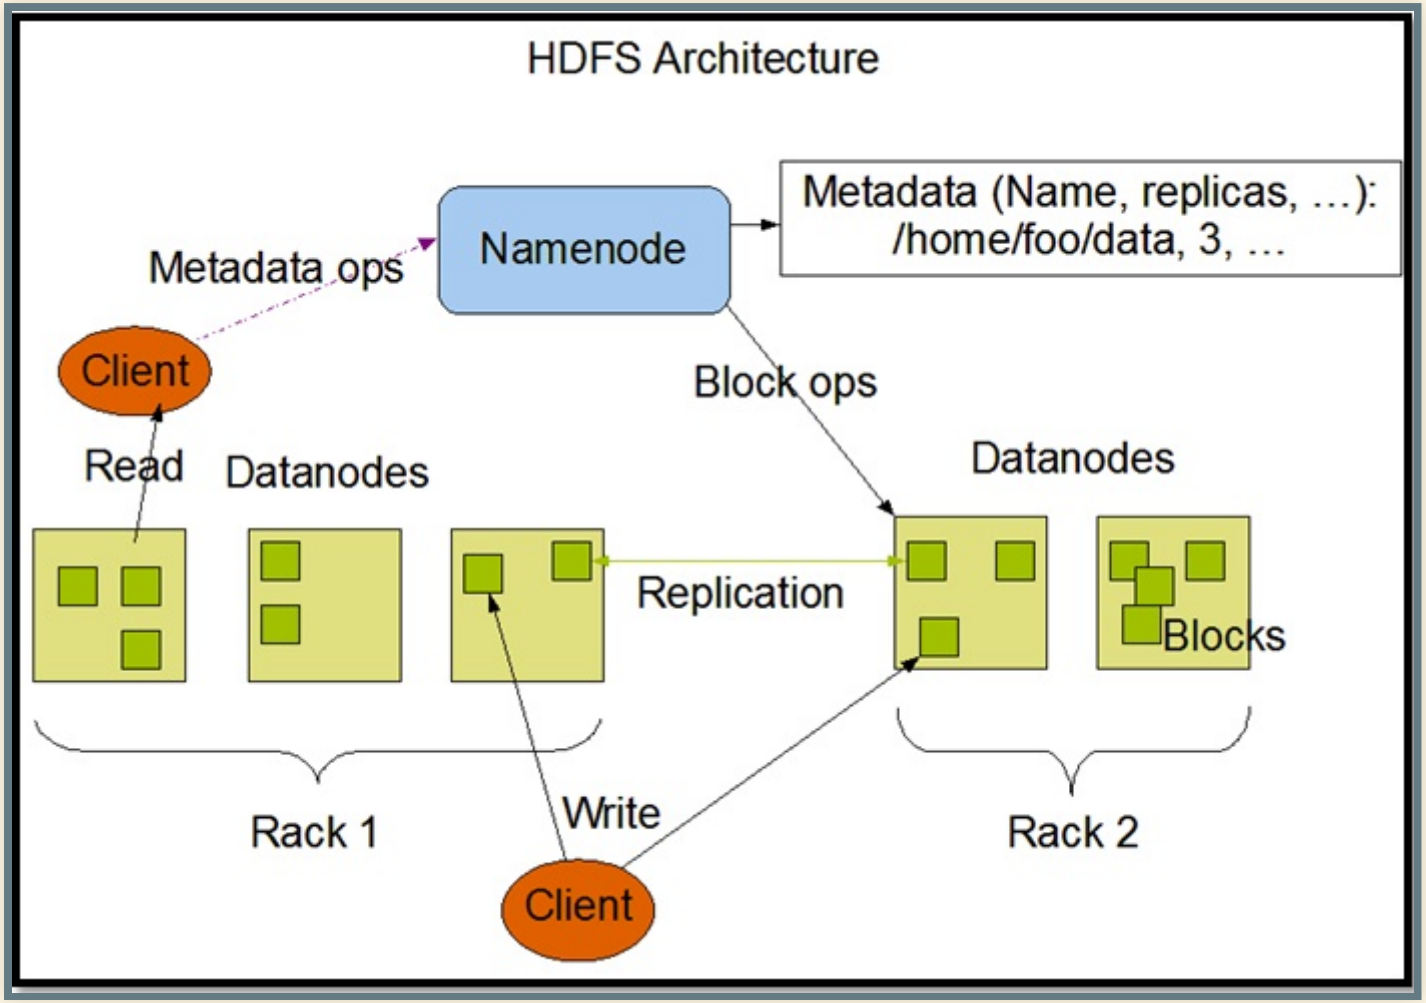

- NN (Namenode) is SPOF (Single Point of Failure)
- Normally HA (High availability) through a standby
- Zookeeper for coordination, bookeeping

WORD COUNT

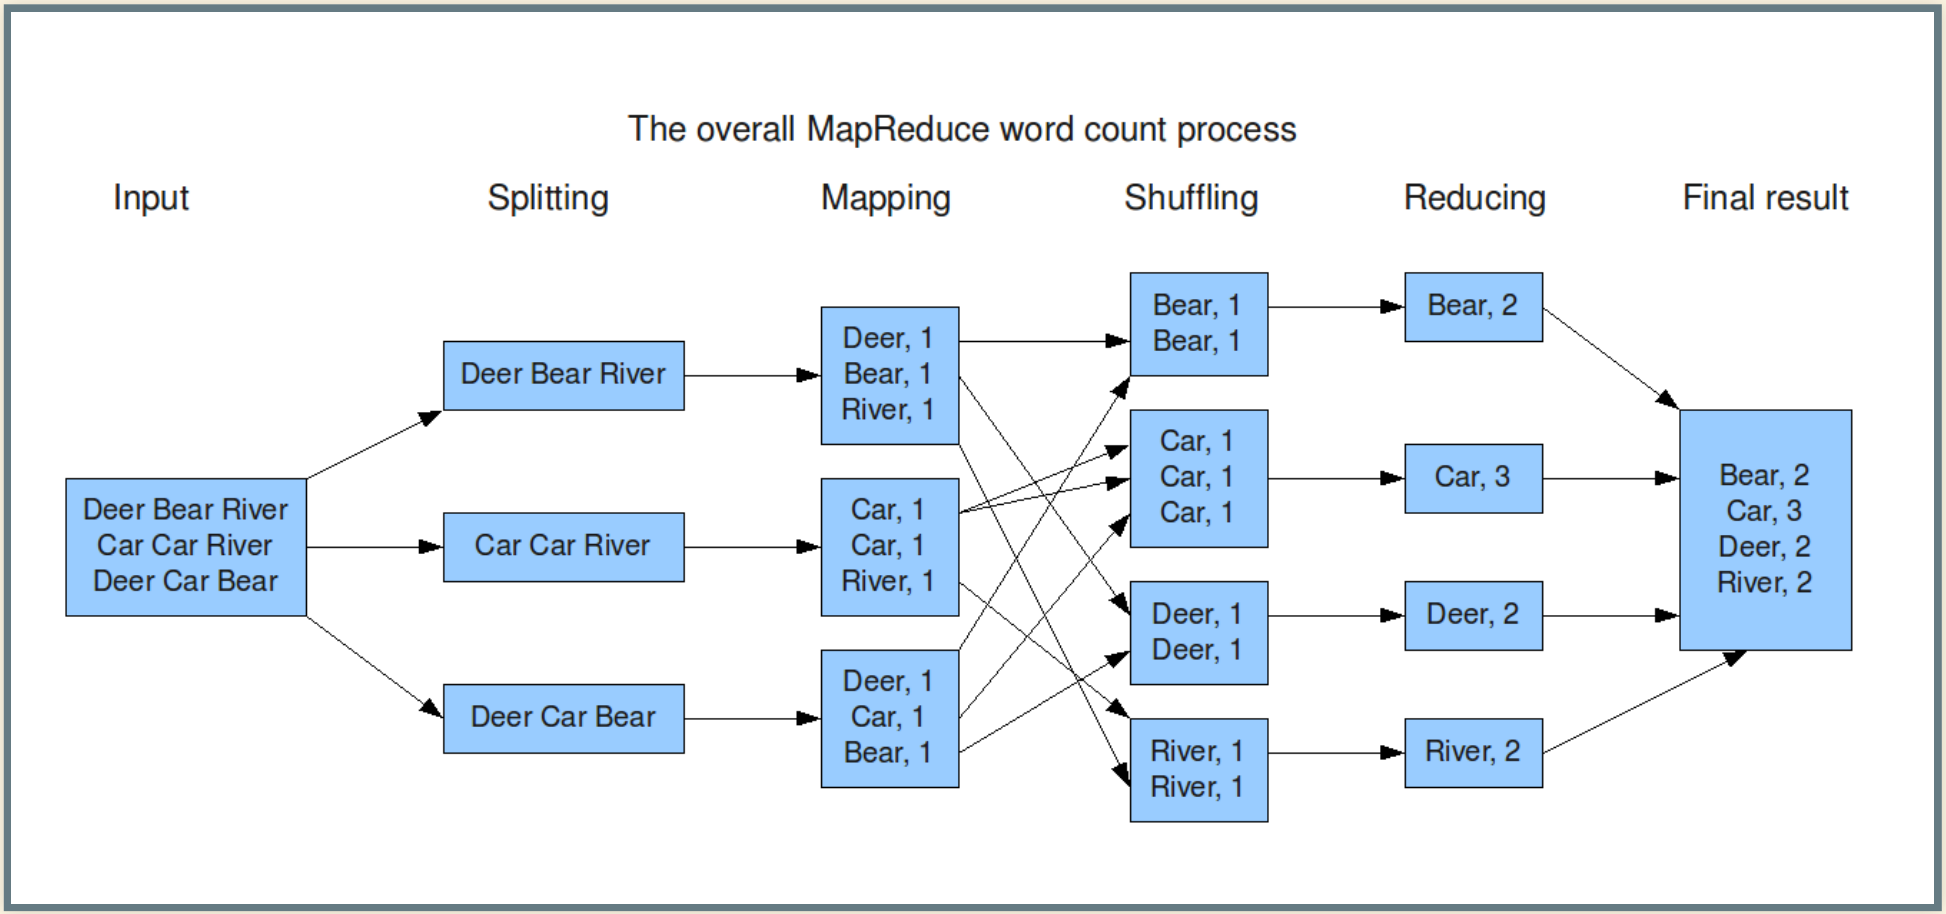
The overall MapReduce word count process
Input
Splitting
Mapping
Shuffling
Reducing
Final result
Deer Bear River
Car Car River
Deer Car Bear
Deer Bear River
Car Car River
Deer Car Bear
Deer, 1
Bear, 1
River, 1
Car, 1
Car, 1
River, 1
Deer, 1
Car, 1
Bear, 1
Bear, 1
Bear, 1
Car, 1
Car, 1
Car, 1
Deer, 1
Deer, 1
River, 1
River, 1
Bear, 2
Car, 3
Deer, 2
River, 2
Bear, 2
Car, 3
Deer, 2
River, 2

MAPREDUCE ARCHITECTURE

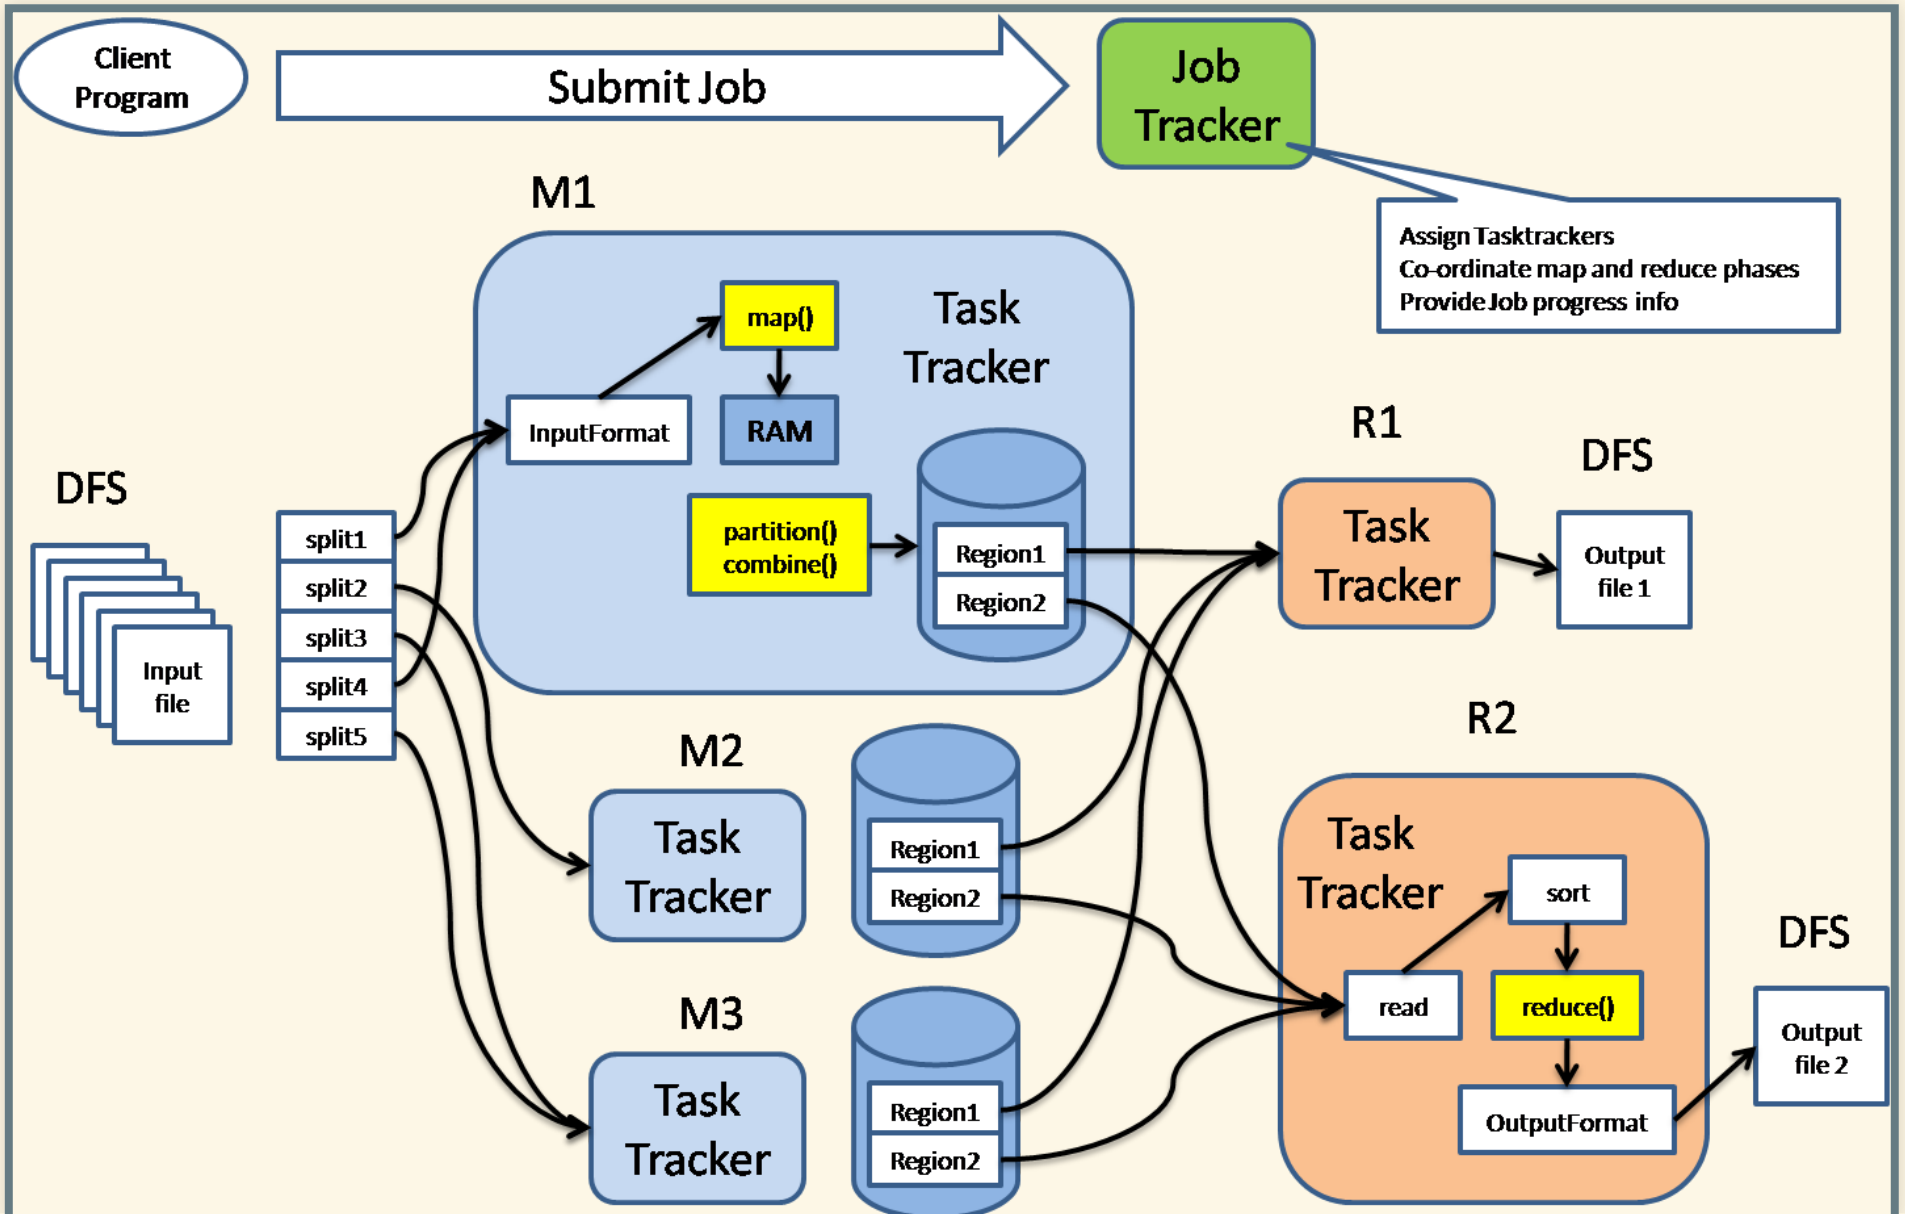

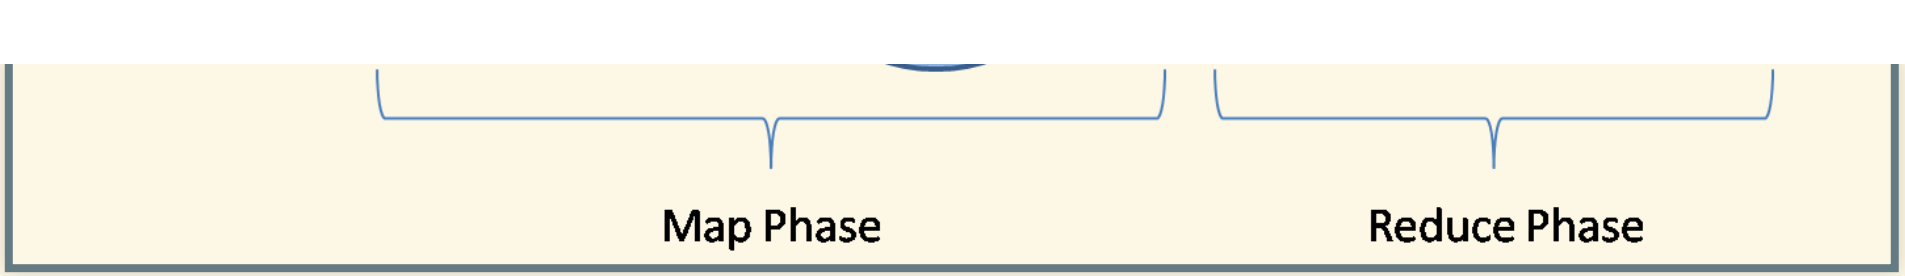
Map Phase
Reduce Phase

QUERYING

- Initially Java

Sample MapReduce (small subset of the entire code which totals nearly 150 lines):

```
public static class MapClass
extends Mapper<WordOffset, Text, Text, IntWritable> {
  private final static String delimiters =
      "',./<>?;:\"[]{}-=_+()&*%^#$!@`~ \\|«»¡¢£¤¥¦©¬®¯±¶·¿";
  private final static IntWritable one = new IntWritable(1);
  private Text word = new Text();
  public void map(WordOffset key, Text value, Context context)
    throws IOException, InterruptedException {
    String line = value.toString();
    StringTokenizer itr = new StringTokenizer(line, delimiters);
    while (itr.hasMoreTokens()) {
       word.set(itr.nextToken());
       context.write(word, one);
    }
  }
}
```

QUERYING

- Apache Pig

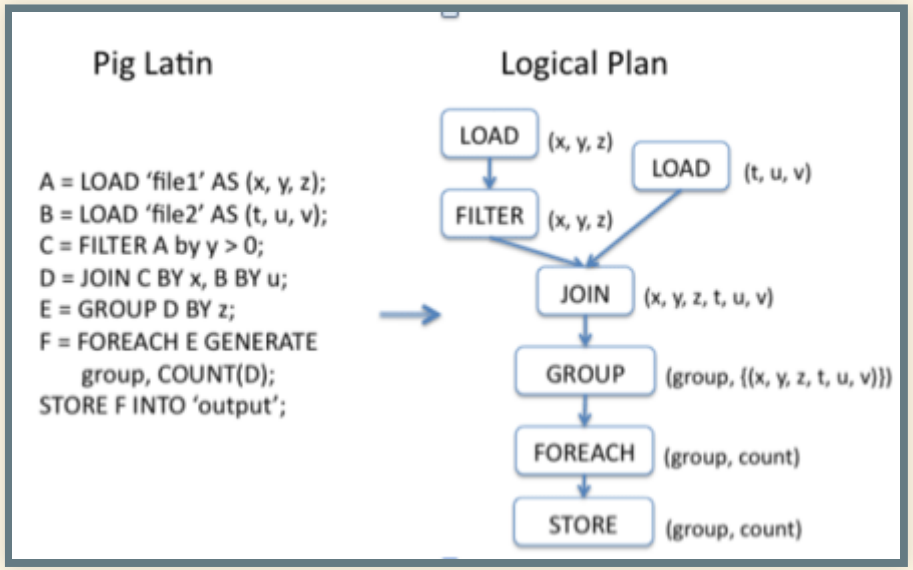

- Apache Hive

```
SELECT pv_users.gender,
count(DISTINCT pv_users.userid),
```

```
count(*), sum(DISTINCT
pv_users.userid) FROM pv_users GROUP
BY pv_users.gender;
```

SQL

- Everyone knows
- Easy for analysts
- But there are actually alot of ways of querying HDFS
    - Cascading, Scalding, Cascalog, etc.

```
class WordCountJob(args : Args) extends Job(args) {
  TypedPipe.from(TextLine(args("input")))
    .flatMap { line => line.split("""\s+""") }
    .groupBy { word => word }
    .size
    .write(TypedTsv(args("output")))
}
```

```
(?- (stdout)
    (<- [?word ?count]
        (sentence :> ?line)
```

```
        (tokenise :< ?line :> ?word)
        (c/count :> ?count)))
```

HDP

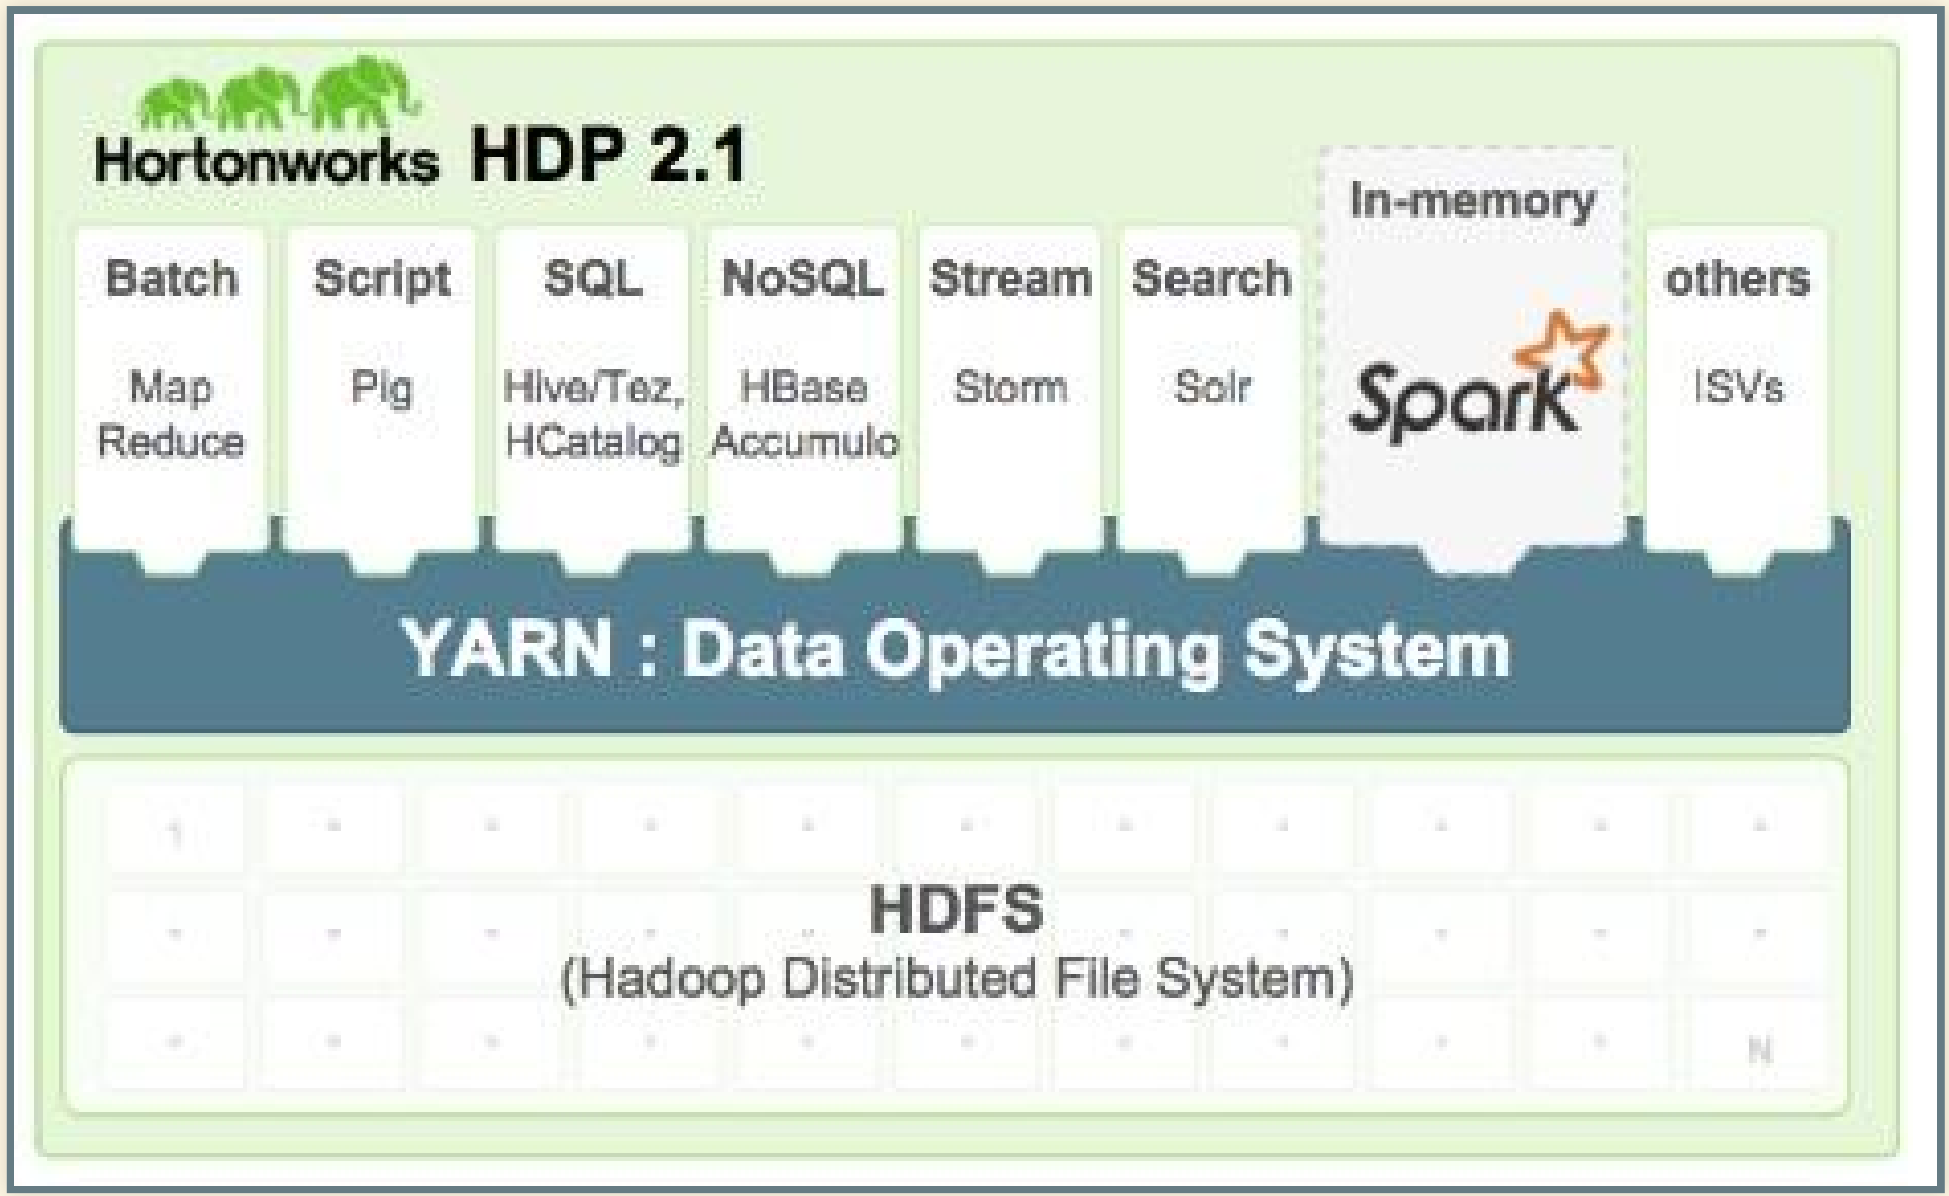
Hortonworks HDP 2.1
Batch
Map Reduce
Script
Pig
SQL
Hive/Tez, HCatalog
NoSQL
HBase Accumulo
Stream
Storm
Search
Solr
In-memory
Spark
others
ISVs
YARN : Data Operating System
HDFS
(Hadoop Distributed File System)

BUSINESS NEEDS

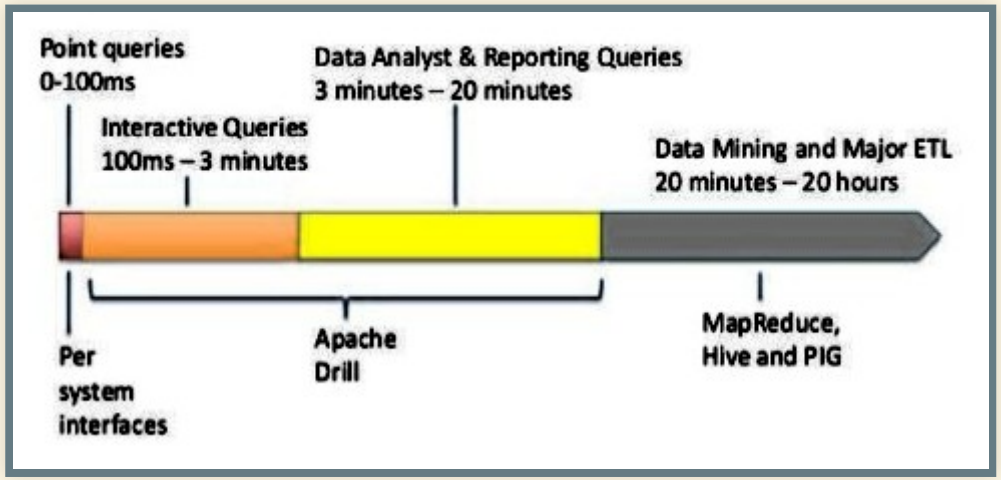

LAMBDA ARCHITECTURE

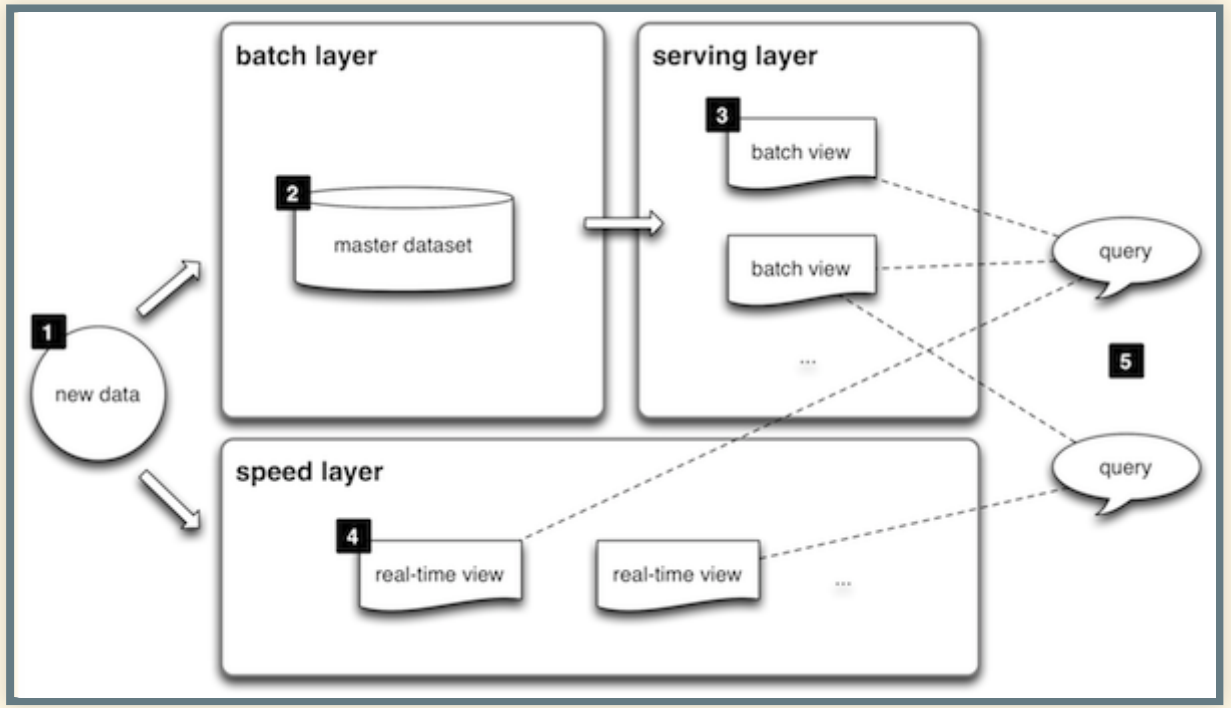

2. PRESTO

ARCHITECTURE

Connectors

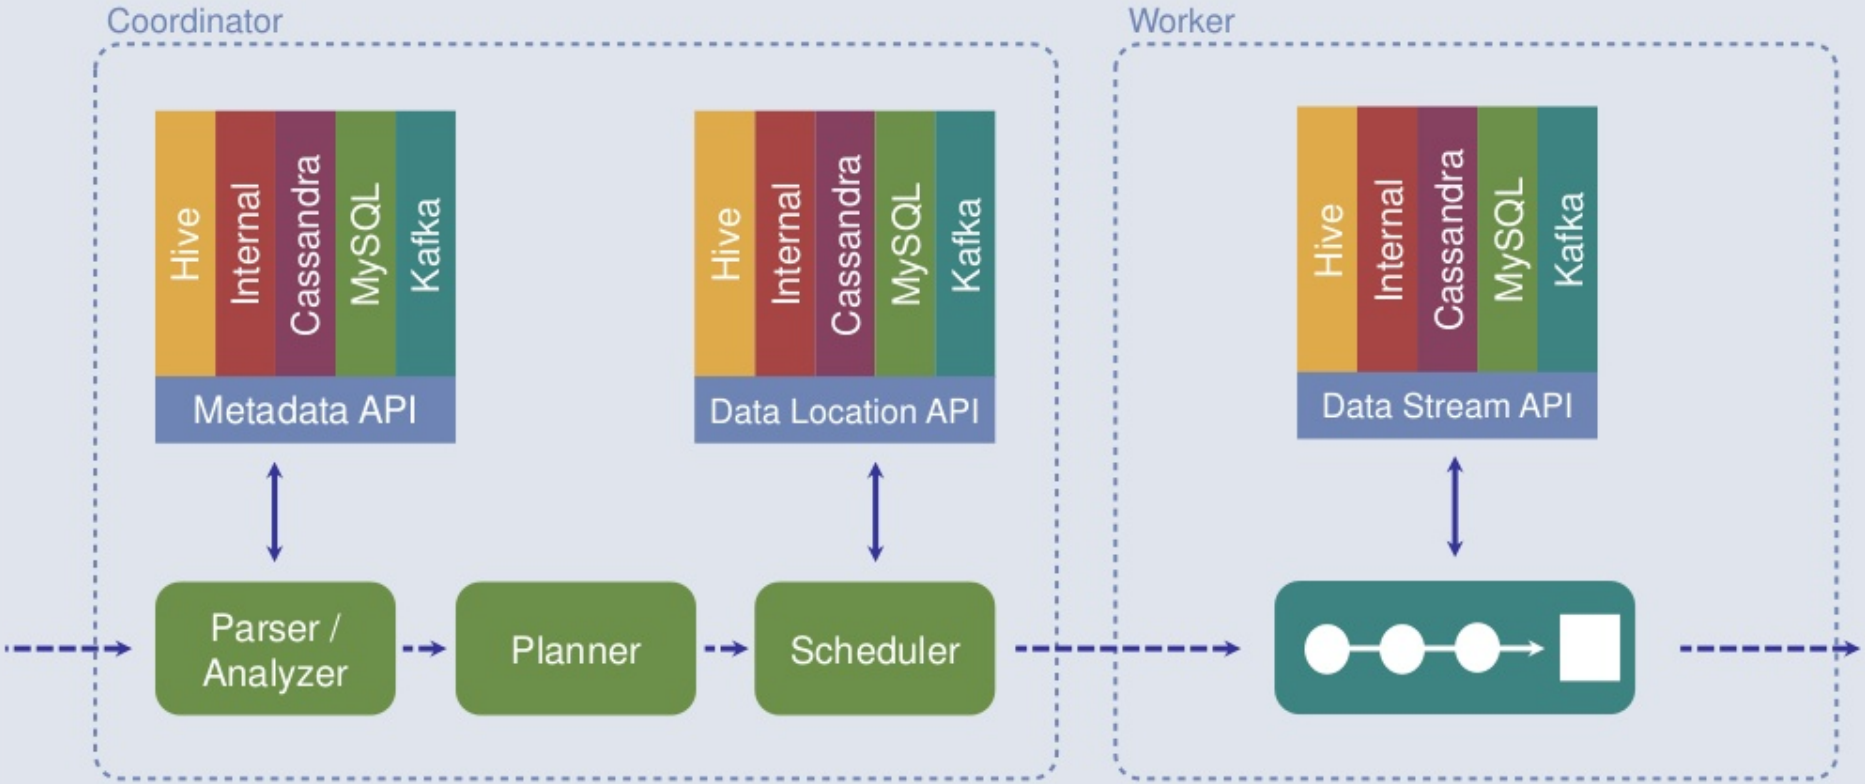
Coordinator
Worker
Hive
Internal
Cassandra
MySQL
Kafka
Metadata API
Hive
Internal
Cassandra
MySQL
Kafka
Data Location API
Hive
Internal
Cassandra
MySQL
Kafka
Data Stream API
Parser / Analyzer
Planner
Scheduler

ARCHITECTURE

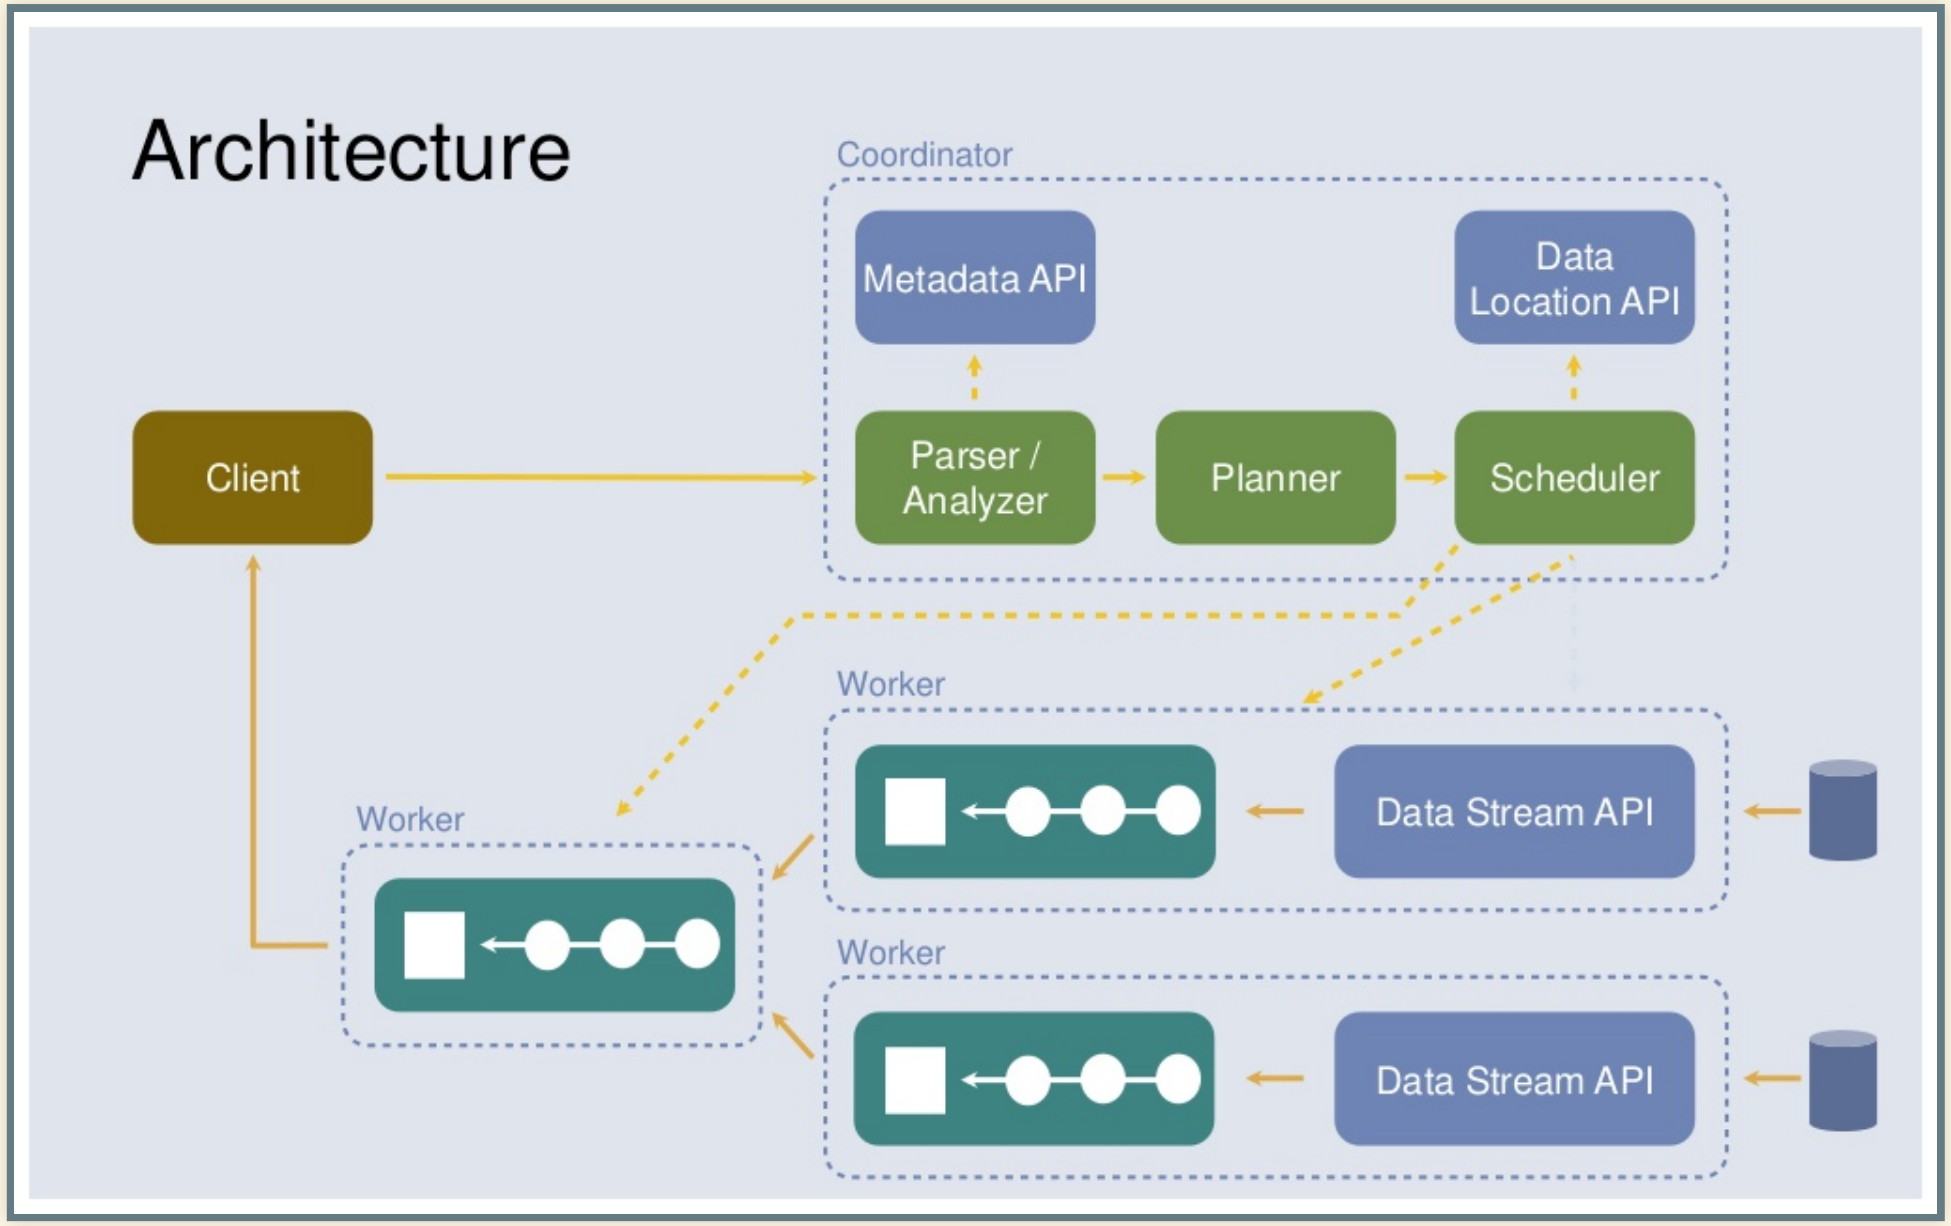
Architecture
Coordinator
Metadata API
Data Location API
Client
Parser / Analyzer
Planner
Scheduler
Worker
Data Stream API
Worker
Worker
Data Stream API

MR VS PRESTO

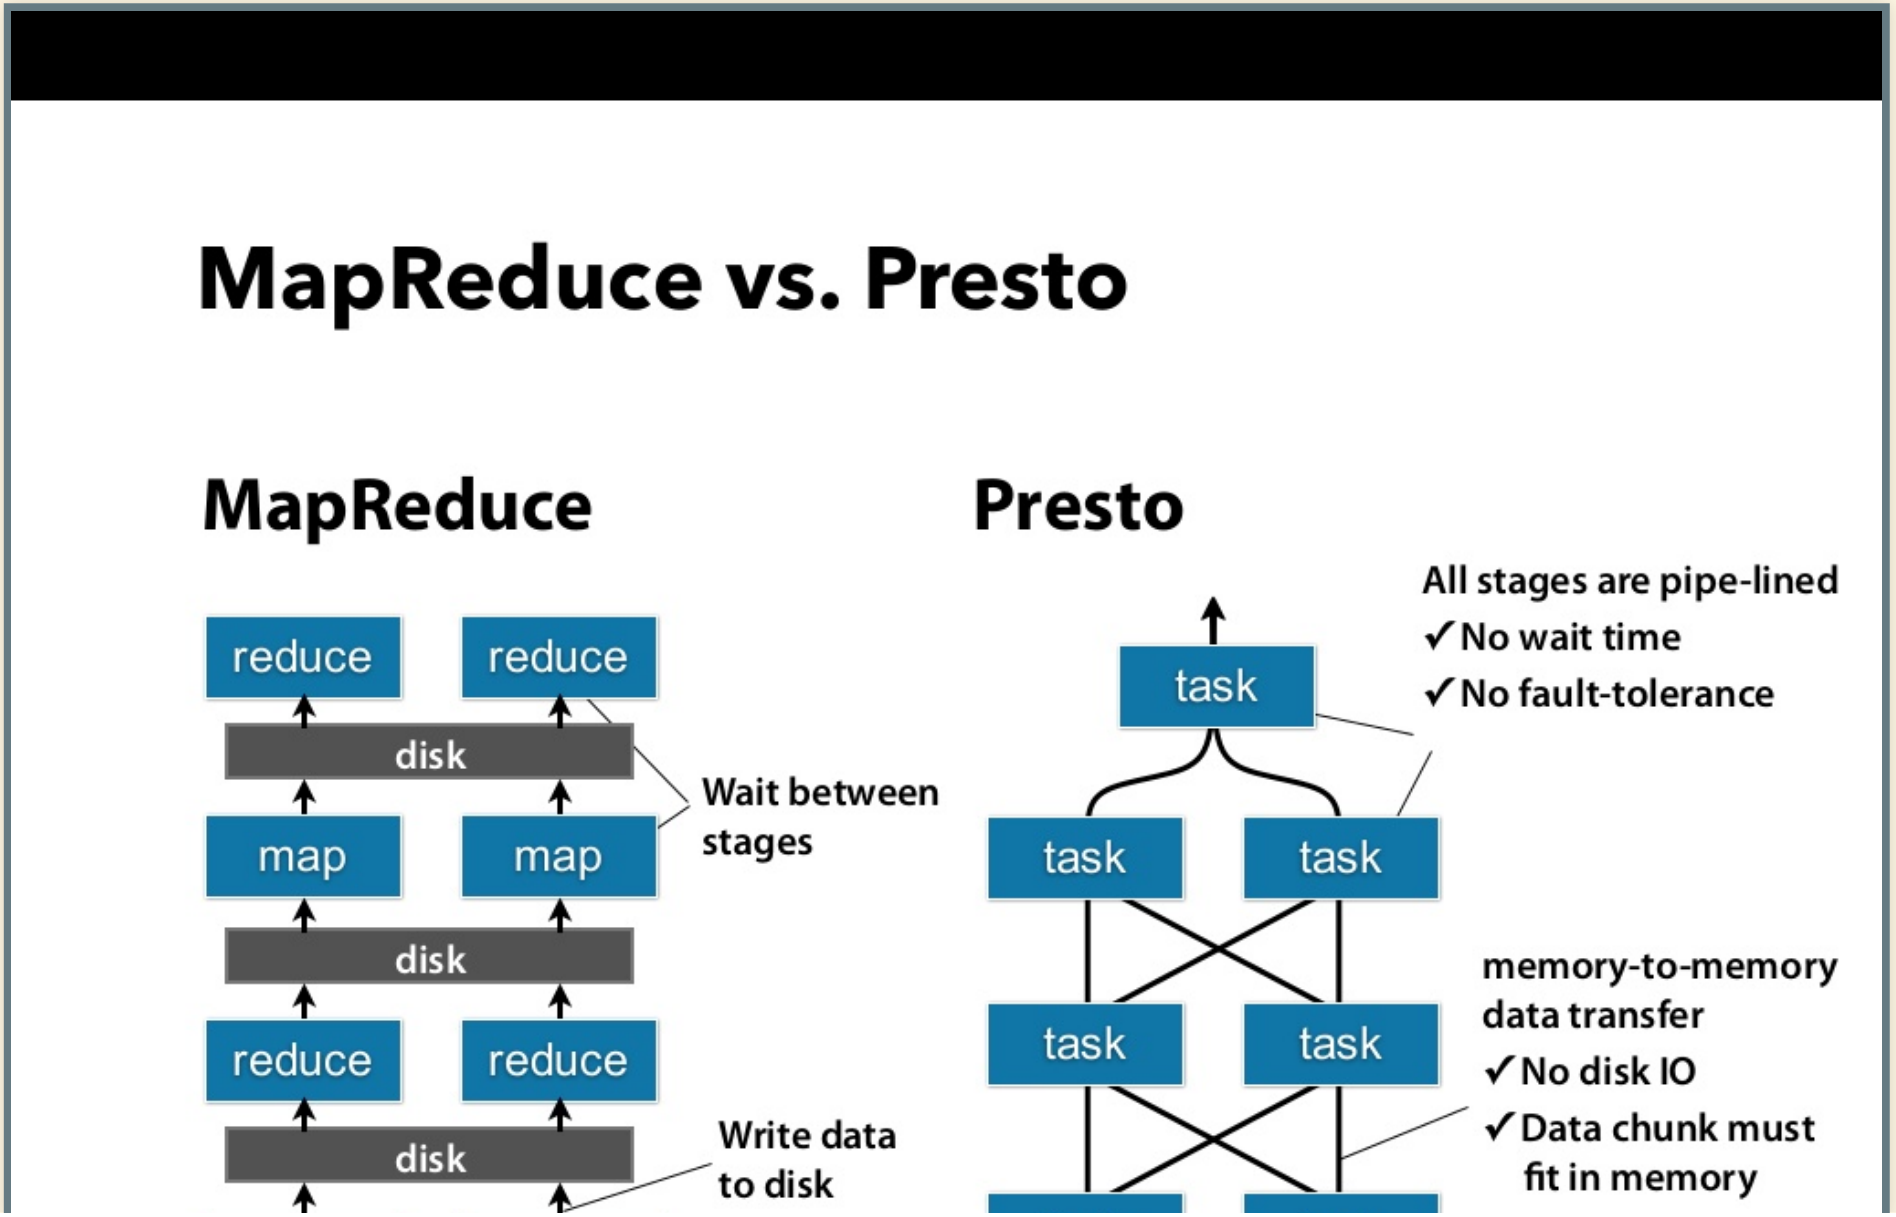
MapReduce vs. Presto
MapReduce
reduce
reduce
disk
map
map
disk
reduce
reduce
disk
Wait between stages
Write data to disk
Presto
All stages are pipe-lined
✓No wait time
✓No fault-tolerance
task
task
task
task
task
memory-to-memory data transfer
✓No disk IO
✓Data chunk must fit in memory


map
map
task
task

DETAILS

- Data must fit in memory
- Uses connectors to several backends
  - Cassandra, Hive, JMX, Kafka, Mysql, Postgres
- Nodes are stateless

SUPPORT

- Ansi SQL
- Approximation functions
- JSON functions
- PrestoML

TRICKS

BEST ONES

- **Vectorized Reader**: read based on column vectors
- **Predicate Pushdown**: use statistics/logic to skip data
- **Lazy Load**: postpone loading until needed
- **Lazy materialization**: postpone decoding until needed

2.1 COLUMNAR STORE

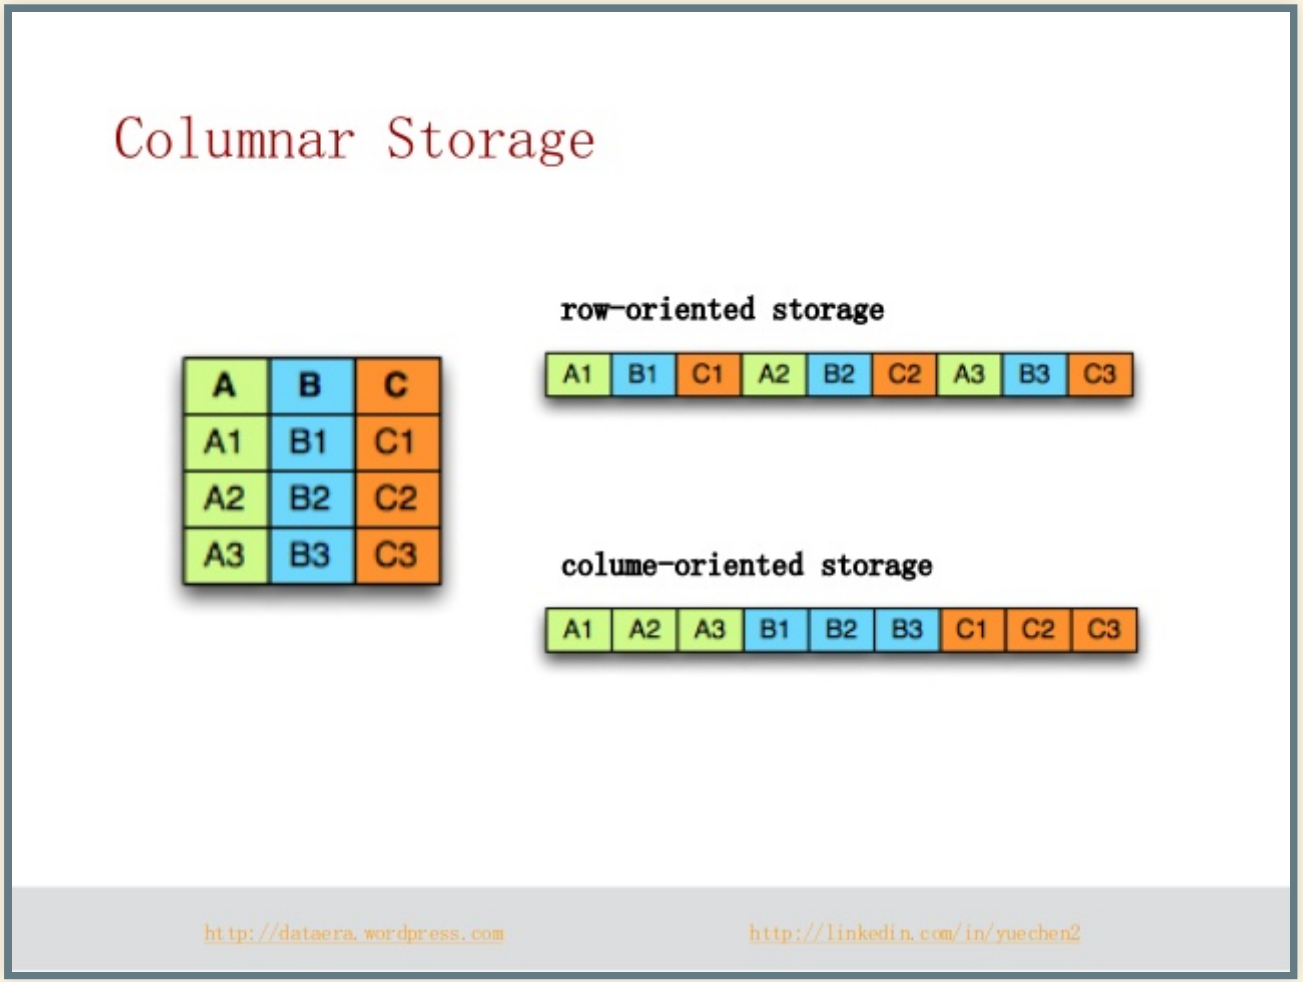

2.2 FILE SIZE COMPARISON

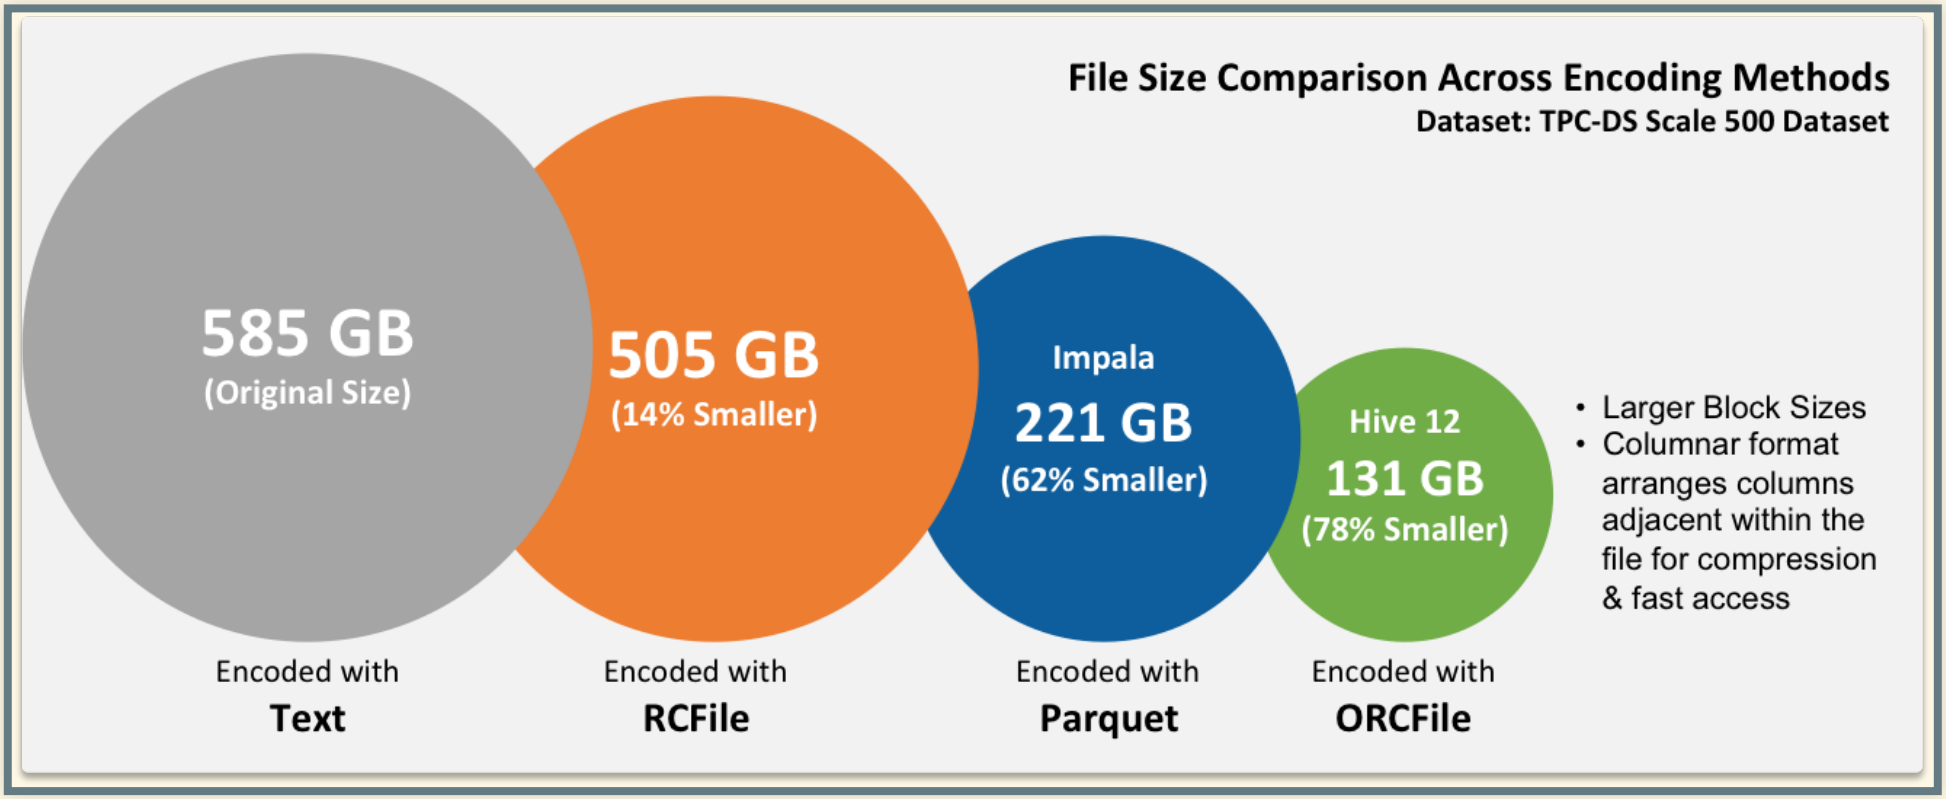

FILE LAYOUT

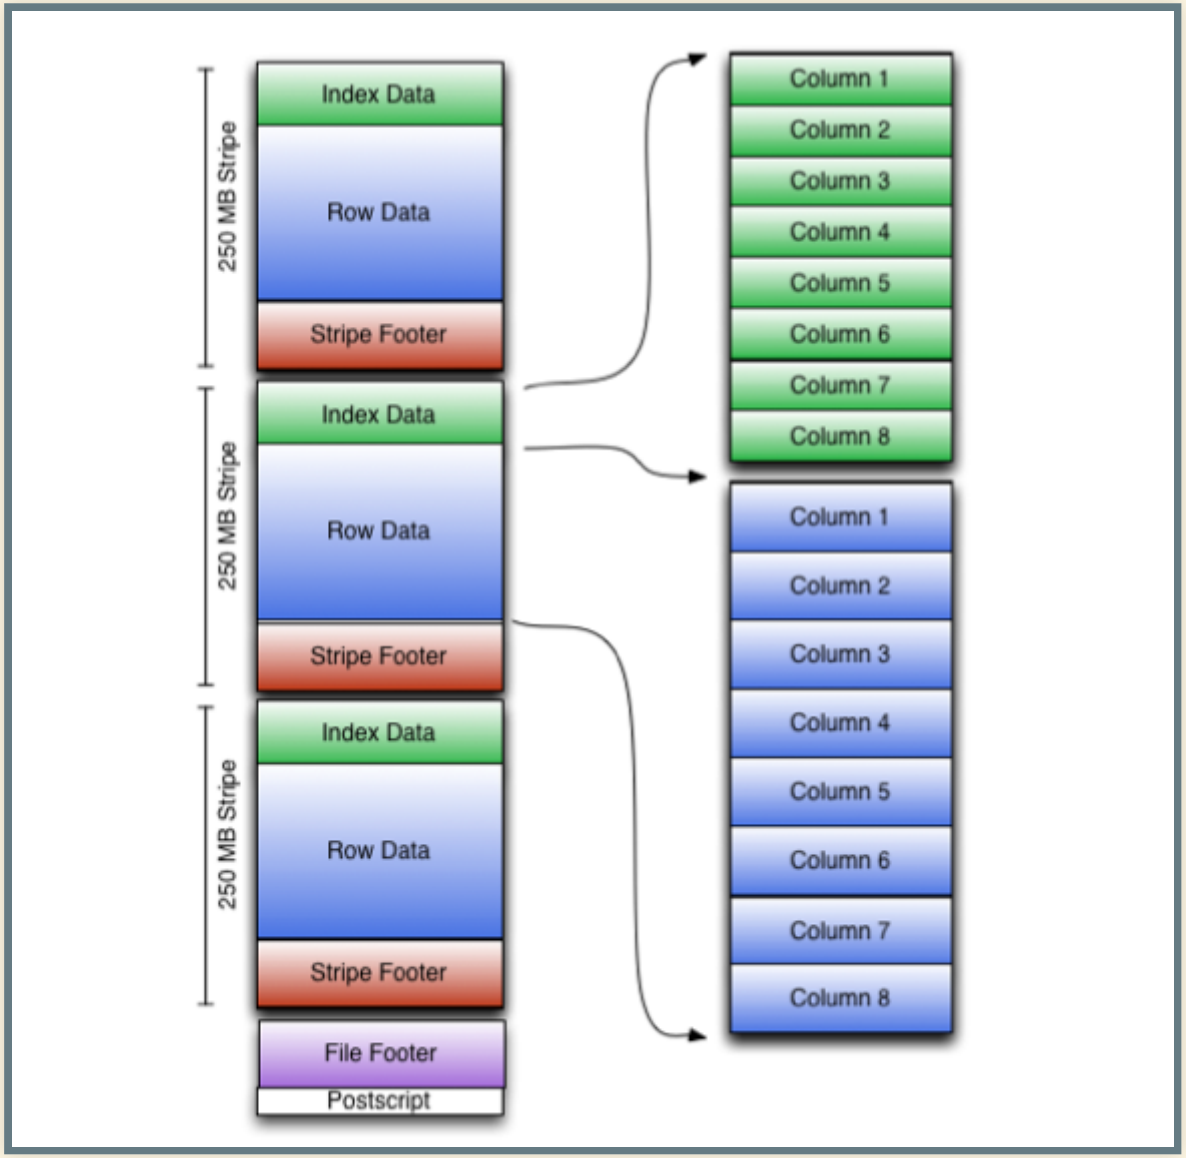
250 MB Stripe
Index Data
Row Data
Stripe Footer
250 MB Stripe
Index Data
Row Data
Stripe Footer
250 MB Stripe
Index Data
Row Data
Stripe Footer
File Footer
Postscript
Column 1
Column 2
Column 3
Column 4
Column 5
Column 6
Column 7
Column 8
Column 1
Column 2
Column 3
Column 4
Column 5
Column 6
Column 7
Column 8

COMPRESSION

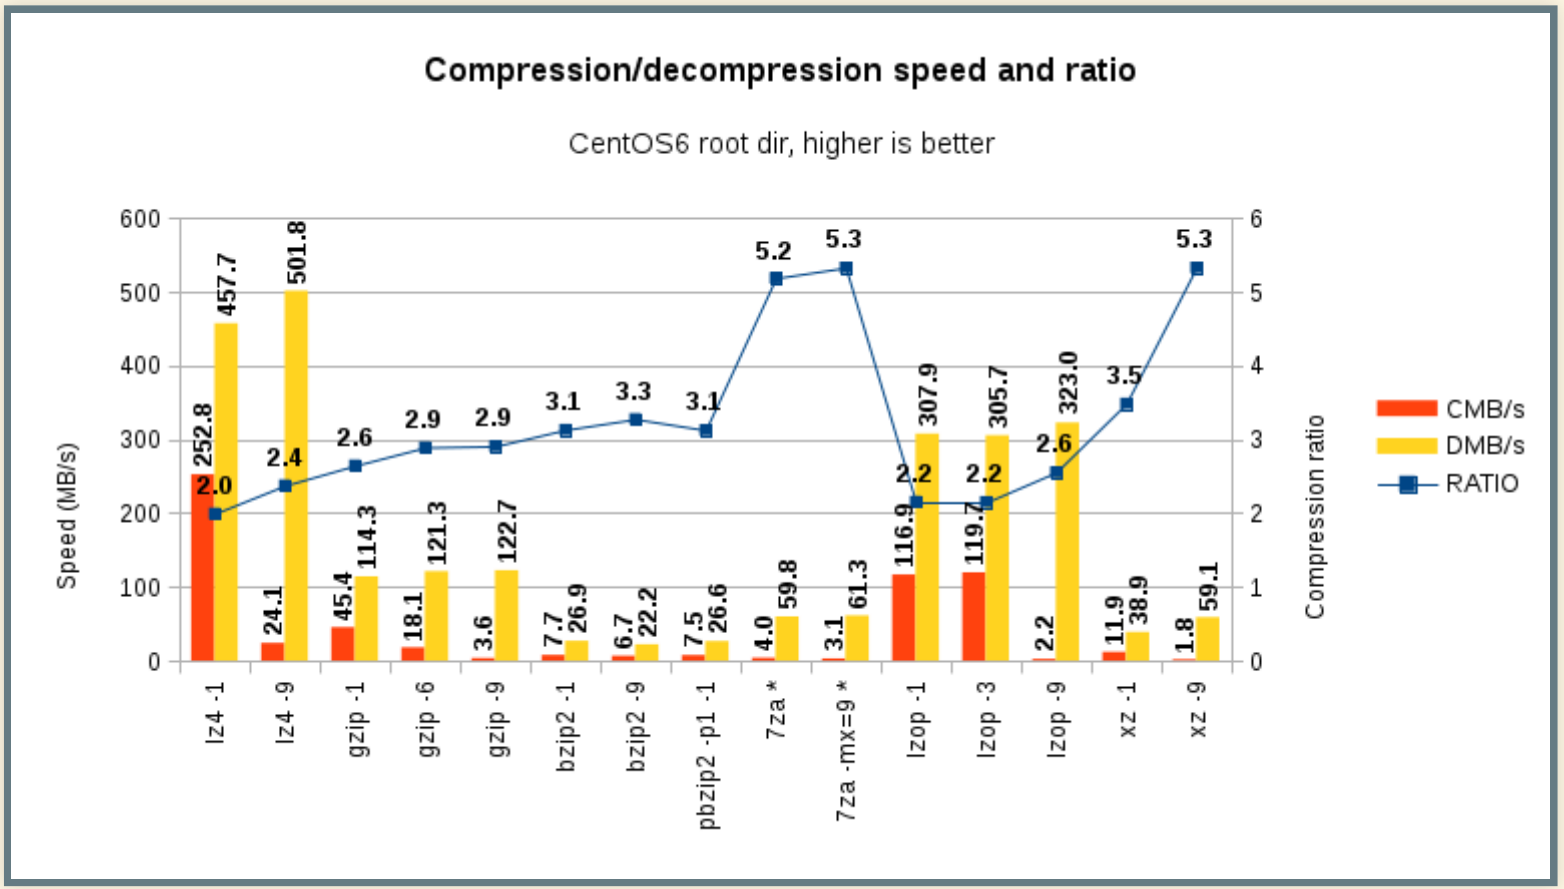

BENCHMARKS

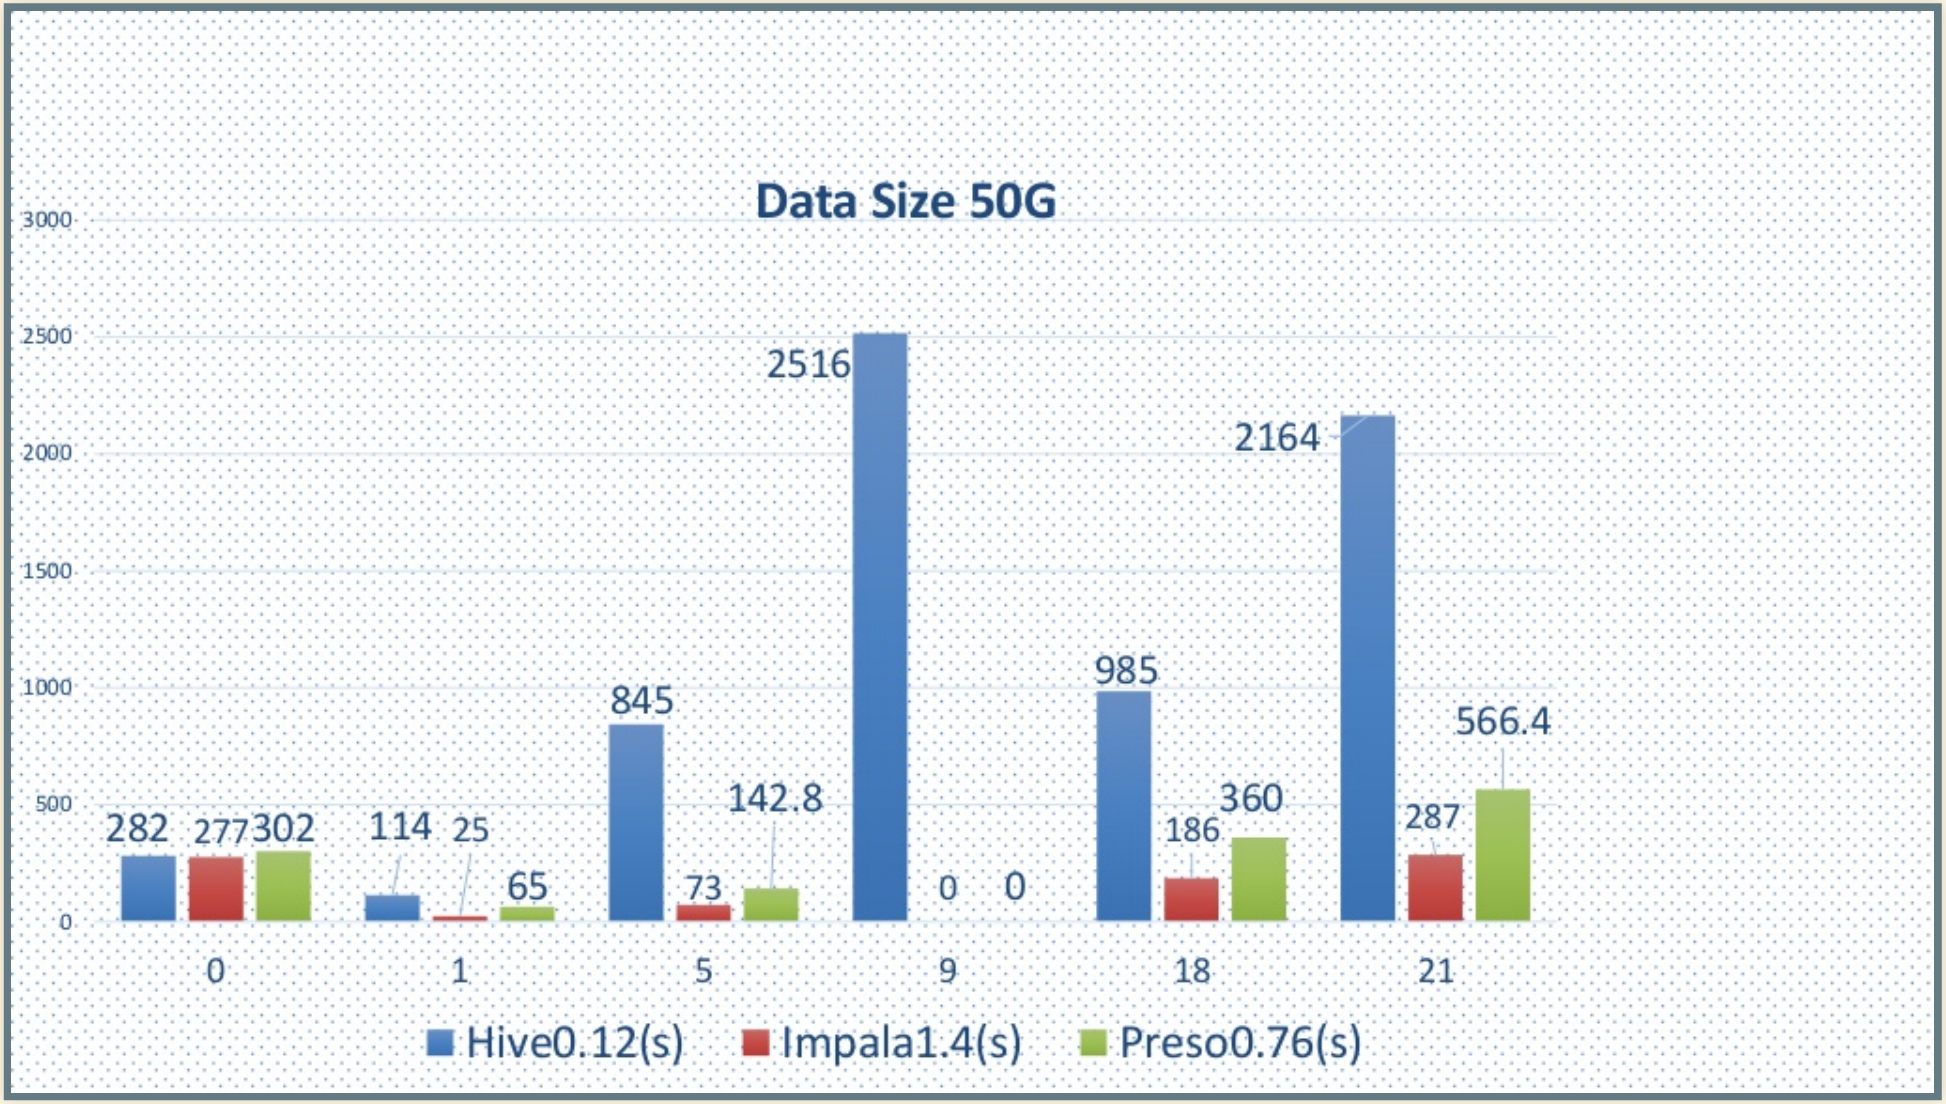
Data Size 50G
3000
2500
2000
1500
1000
500
0
282
277
302
114
25
65
845
73
142.8
2516
0
0
985
186
360
2164
287
566.4
0
1
5
9
18
21
Hive0.12(s)
Impala1.4(s)
Preso0.76(s)

BENCHMARKS

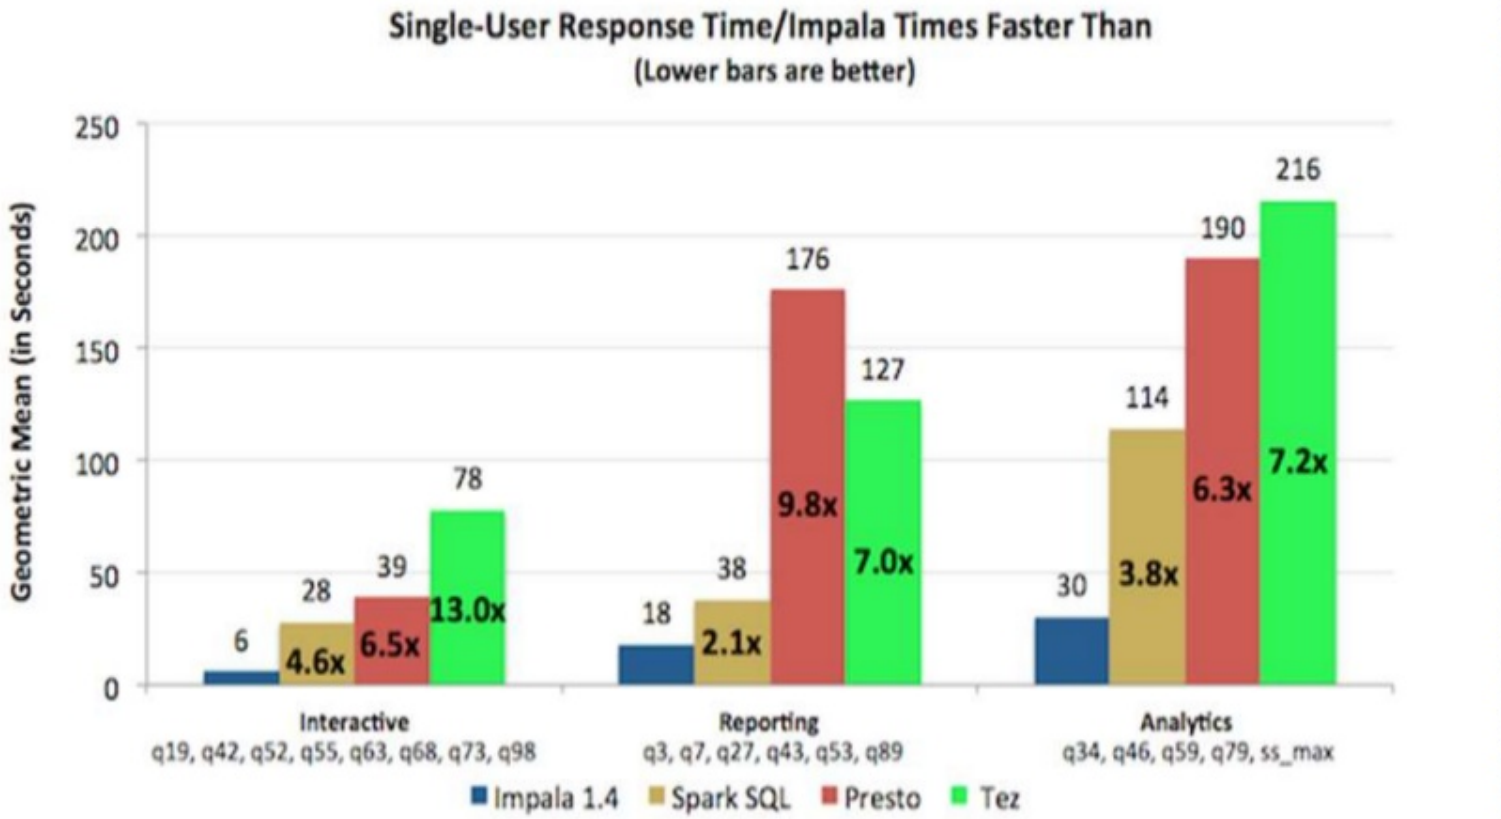
Single-User Response Time/Impala Times Faster Than
(Lower bars are better)
Geometric Mean (in Seconds)
250
200
150
100
50
0
6
28
4.6x
39
6.5x
78
13.0x
18
38
2.1x
176
9.8x
127
7.0x
30
114
3.8x
190
6.3x
216
7.2x
Interactive
q19, q42, q52, q55, q63, q68, q73, q98
Reporting
q3, q7, q27, q43, q53, q89
Analytics
q34, q46, q59, q79, ss_max
Impala 1.4
Spark SQL
Presto
Tez

OTHERS

- Fast inner loops (related to CPU cache)
- sun.misc.Unsafe (non-gc memory access)
- Pipelining, streaming

NOTES

- SQL parser is being used in other projects (ex: crate.io)
- Airpal - UI for presto
- There's also Impala, Apache Drill, Apache Tajo, Redshift, Spark, etc
- PaaS - Presto as a service called Qubole

NETFLIX

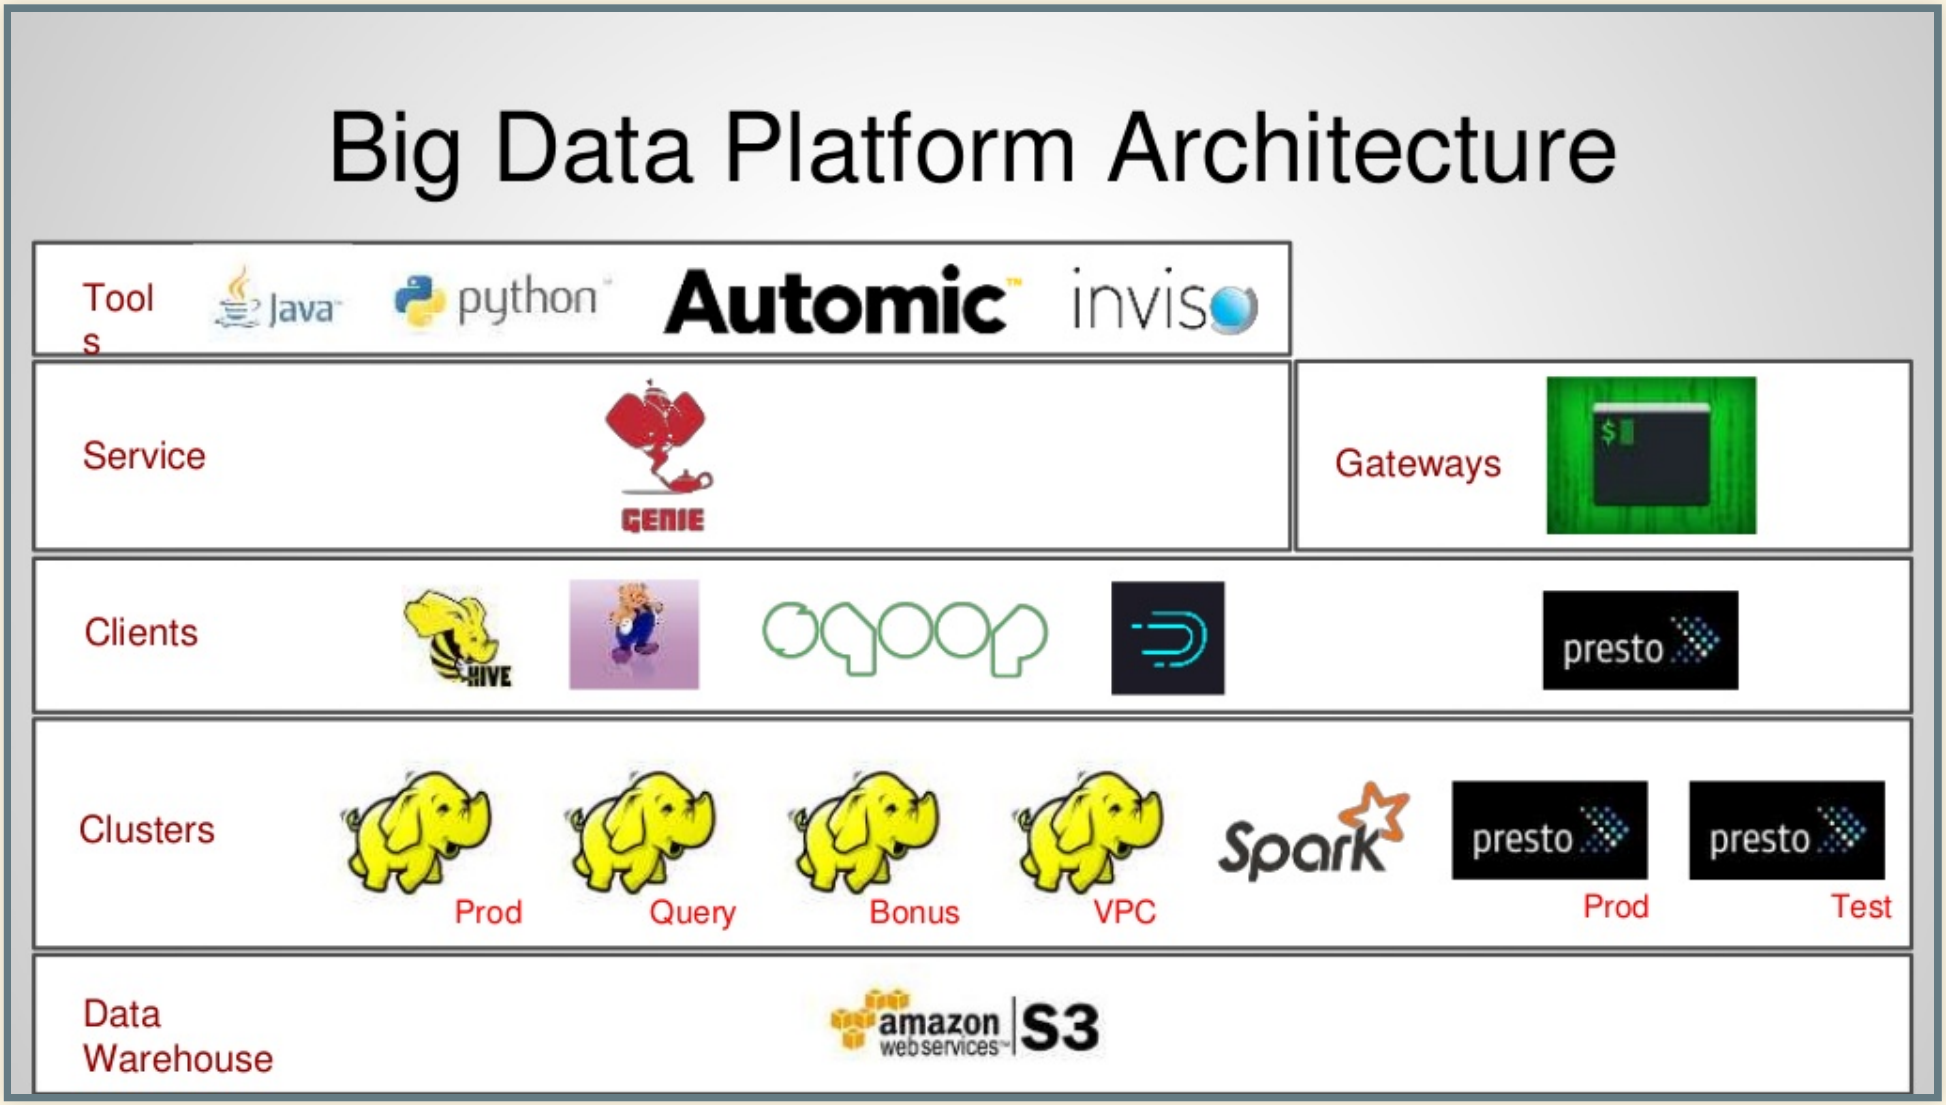
Big Data Platform Architecture
Tools
Java
python
Automic
invisoglobal
Service
GENIE
Gateways
Clients
HIVE
presto
Clusters
Prod
Query
Bonus
VPC
Spark
presto
Prod
presto
Test
Data Warehouse
amazon web services
S3

http://www.slideshare.net/treasure-data/2015-0311td-techtalkinternalsofprestoservice?qid=702c79ef-0632-476b-abb0-0aaff121cf00&v=default&b=&from_search=12